

Task-4:

HEART DISEASE PREDICTION USING LOGISTIC REGRESSION

Submitted By

Name: A D S Abhishek

Internship: CodeAlpha – Machine Learning Internship

Table of Contents

1. **Abstract**
2. **Introduction**
3. **Problem Statement**
4. **Objectives**
5. **Dataset Description**
6. **Data Preprocessing**
7. **Exploratory Data Analysis (EDA)**
8. **Methodology**
9. **Model Building and Prediction**
10. **Results and Evaluation**
11. **Sample Predictions**

12. Conclusion

13. References

Abstract

Heart disease is one of the leading causes of death worldwide, making early detection and prediction extremely important. This project aims to develop a machine learning-based system that predicts the possibility of heart disease using patient medical information such as age, sex, chest pain type, blood pressure, cholesterol level, maximum heart rate, and other clinical parameters.

The Heart Disease dataset containing 1025 patient records and 14 attributes was used. Logistic Regression, a supervised machine learning classification algorithm, was applied to predict the presence or absence of heart disease.

The model was trained and evaluated using a train-test split approach. Its performance was measured using classification metrics such as accuracy, precision, recall, and F1-score. The developed system demonstrates how machine learning techniques can be effectively used in healthcare applications to support early disease prediction and assist in clinical decision-making.

Introduction

Heart disease is a major health concern worldwide and is responsible for a large number of deaths every year. Early prediction of heart disease can help healthcare professionals provide timely treatment and preventive measures. Machine learning techniques can analyze medical data and identify patterns that are difficult to observe manually.

This project develops a Heart Disease Prediction System using Logistic Regression to classify whether a patient is likely to have heart disease based on medical parameters.

Problem Statement

To develop a machine learning-based system that predicts whether a patient is likely to have heart disease or not based on medical attributes such as age, sex, chest pain type, blood pressure, cholesterol level, maximum heart rate, and other clinical parameters. The goal is to assist in early diagnosis and support healthcare professionals in making informed decisions.

Objective: Predict the possibility of diseases based on patient data.

- **To analyze the Heart Disease dataset.**
- **To preprocess and visualize medical data.**
- **To build a Logistic Regression model for disease prediction.**
- **To evaluate model performance using classification metrics.**

Approach: Apply classification techniques to structured medical datasets.

Key Features:

- **Use features like symptoms, age, blood test results, etc.**
- **Algorithms Used: Logistic Regression, Random Forest**
- **Datasets: Heart disease, Diabetes, Breast Cancer (UCI ML Repository).**

Dataset Description

- **Dataset: Heart Disease Dataset (Kaggle/UCI Repository)**
- **Rows (Instances): 1025**
Columns (Attributes): 14
Input Features: 13
Target Variable: target
- **Target:**
0 → No Heart Disease
1 → Presence of Heart Disease
- **Features:**
age, sex, cp, trestbps, chol, fbs, restecg, thalach, exang, oldpeak, slope, ca, thal.
- **Missing Values: 0**
Data Types: 13 Integer columns and 1 Float column.

Attribute	Description
age	Patient's age (in years)
sex	Gender (0 = Female, 1 = Male)
cp	Type of chest pain
trestbps	Resting blood pressure
chol	Cholesterol level
fbs	Fasting blood sugar (1 = High, 0 = Normal)
restecg	Resting ECG results
thalach	Maximum heart rate achieved
exang	Exercise-induced chest pain (1 = Yes, 0 = No)
oldpeak	ST depression caused by exercise
slope	Slope of the exercise ST segment
ca	Number of major blood vessels

thal
target

Thalassemia type
Heart disease prediction (0 = No Disease, 1 = Disease)

Data Preprocessing

Data preprocessing is an important step in machine learning that involves preparing the dataset for model training and improving data quality.

- **Data inspection using:**

df.head()

df.info()

df.describe()

df.isnull().sum()

- **Missing values: None**
- **Duplicate values: None**
- **Encoding required: No**
- **All features were already numerical and ready for model training.**

Therefore, the dataset was directly divided into input features (X) and target variable (y) and used for training the Logistic Regression model.

```
Disease_prediction.py 1 X
Disease_prediction.py > ...
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sn
5 from sklearn.model_selection import train_test_split
6 from sklearn.metrics import precision_score, recall_score, f1_score
7 from sklearn.linear_model import LogisticRegression
8 from sklearn.metrics import accuracy_score
9 from sklearn.metrics import classification_report
```

These are the required and imported libraries/packages.

Dataset link used from the Kaggle resource: and required link.

<https://www.kaggle.com/datasets/johnsmith88/heart-disease-dataset>

Exploratory Data Analysis (EDA)

1. Age Distribution Histogram
2. Heart Disease Distribution Count Plot
3. Correlation Heatmap
4. Confusion Matrix Heatmap

Methodology:

Heart Dataset



Data Preprocessing

↓
Exploratory Data Analysis

↓
Feature Selection

↓
Train-Test Split

↓
Logistic Regression Model

↓
Prediction


↓
Model Evaluation

↓
Conclusion



Screenshots of the Data and the executed code:

```
df = pd.read_csv("D:/heart.csv")
```



The screenshot shows a code editor with a file named `Disease_prediction.py`. The script contains the following code:

```
12 # Data Preprocessing
13 print(df.head())
14 print(df.info())
15 print(df.isnull().sum())
16 print(df.describe())
17
```


The terminal output shows the execution of the script. It displays the first five rows of the dataset, the output of `df.info()` (1025 entries, 14 columns), and the output of `df.isnull().sum()` (all zeros). The output of `df.describe()` is also shown, providing summary statistics for each column.

#	Column	Non-Null count	Dtype
0	age	1025 non-null	int64
1	sex	1025 non-null	int64
2	cp	1025 non-null	int64
3	trestbps	1025 non-null	int64
4	chol	1025 non-null	int64
5	fbs	1025 non-null	int64
6	restecg	1025 non-null	int64
7	thalach	1025 non-null	int64
8	exang	1025 non-null	int64
9	oldpeak	1025 non-null	float64
10	slope	1025 non-null	int64
11	ca	1025 non-null	int64
12	thal	1025 non-null	int64
13	target	1025 non-null	int64

dtypes: float64(1), int64(13)
memory usage: 112.2 KB

```
print(df.isnull().sum())
```

```
print(df.describe())
```



Showing the Histogram:

```
plt.figure(figsize=(8,5))
```

```
plt.hist(df['age'], bins=10)
```


```
plt.xlabel("Age")
```



```
plt.ylabel("Number of Patients")
```

```
plt.title("Age Distribution of Patients")
```

```
plt.show()
```



using the Matplotlib: Disease vs No Disease (Count Plot)

Purpose: Shows the number of patients with and without heart disease.

```
plt.figure(figsize=(6,5))
```

```
df['target'].value_counts().plot(kind='bar')
```


```
plt.xlabel("Heart Disease")
```

```
plt.ylabel("Count")
```

```
plt.title("Heart Disease Distribution")
```

```
plt.xticks([0,1], ['No Disease','Disease'])
```

```
plt.show().
```



Heatmap Representation:

`correlation = df.corr()`

`plt.figure(figsize=(12,8))`

`sn.heatmap(`

`correlation,`

`annot=True,` *# Shows correlation values inside each cell*


`cmap="coolwarm",` *# Color scheme*

`fmt=".2f"` *# Shows values up to 2 decimal places*

`)`

```
plt.title("Correlation Heatmap of Heart Disease Dataset") plt.show().
```

Figure 1



Feature Selection

```
X = df.drop('target',axis=1)
y = df['target']
```

X = input features

y = output label.

Train-Test Split

```
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=42)
```

Test_size = 0.2 means

80% → Training.

20% → Testing. And random_state=42 that means we are having the Standard dataset for the Model evaluation and prediction.

Model Building & Prediction

Algorithm Used: Logistic Regression

Train-Test Split:

Training Data = 80%

Testing Data = 20%

Random State = 42

Feature Matrix: X = All medical attributes

Target Variable: y = target

```
model = LogisticRegression(max_iter = 1000)
model.fit(X_train,y_train)

y_pred = model.predict(X_test)
```

We are taking the Logistic Regression() for the model evaluation from the mentioned Model buildings.

Sample:

```
print(y_test.values[:30])
print(y_pred[:30])
```

```
56 y_pred = model.predict(X_test)
57
58 print("Values of y_test values ",y_test.values[:30])
59 print("The predicted values: ",y_pred[:30])
60
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

```
PS E:\Project_Code_Alpha> & "C:\Users\ADSS ABHISHEK\AppData\Local\Programs\Python\Python312\python.exe" e:/Project_Code_Alpha/Disease_prediction.py
Values of y_test values [1 1 0 1 0 1 0 0 1 0 1 0 1 1 0 0 0 1 1 0 1 0 0 0 1 1 1 1 0 0]
The predicted values: [1 1 0 1 0 1 0 0 1 0 1 0 1 1 0 1 0 1 1 0 1 0 0 0 1 1 1 1 0 1]
```

from this output we can clearly see that there are some mismatch values from which the actual values to the predicted values: this shows we need to check the Model's Accuracy for the clear Picture.

```
61 print(accuracy_score(y_test, y_pred))
62
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

```
PS E:\Project_Code_Alpha> & "C:\Users\ADSS ABHISHEK\AppData\Local\Programs\Python\Python312\python.exe" e:/Project_Code_Alpha/Disease_prediction.py
0.7951219512195122
```

So our Model got the score of the 79.5% accuracy for the prediction of the Data.

Confusion_matrix():


Disease_prediction.py X

Disease_prediction.py > ...

```
63 from sklearn.metrics import confusion_matrix
64 cm = confusion_matrix(y_test,y_pred)
65 print(cm)
66
67 plt.figure(figsize=(7,5)) # this part of the code is for the printing of the cm graph clearly:
68 sn.heatmap(cm, annot=True)
69 plt.xlabel('Predicted')
70 plt.ylabel('Truth')
71 plt.show()
72
73 print(classification_report(y_test,y_pred))
74
75 sample = [[63,1,3,145,233,1,0,150,0,2.3,0,0,1]]
76 print(model.predict(sample))
77
78
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

```
PS E:\Project_Code_Alpha> & "C:\Users\ADSS ABHISHEK\AppData\Local\Programs\Python\Python312\python.exe" e:/Project_Code_Alpha/Disease_prediction.py
0.7951219512195122
[[73 29]
 [13 90]]
```



Classification Report(): we can clearly visualize the values of the Precision, accuracy, F1-score

```
72
73 print(classification_report(y_test,y_pred))
74
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

```
PS E:\Project_Code_Alpha> & "C:\Users\ADSS ABHISHEK\AppData\Local\Programs\Python\Python312\python.exe" e:/Project_Code_Alpha/Disease_prediction.py
```

	precision	recall	f1-score	support
0	0.85	0.72	0.78	102
1	0.76	0.87	0.81	103
accuracy			0.80	205
macro avg	0.80	0.79	0.79	205
weighted avg	0.80	0.80	0.79	205

Results and Evaluation

Metrics Used:

- Accuracy Score
- Confusion Matrix
- Precision
- Recall
- F1-Score

Model Accuracy:

79.5%

The Logistic Regression model achieved an accuracy of approximately 79.5% on the test dataset, indicating that the model can effectively classify patients with and without heart disease.

Prediction of Samples:

Sample1: Prediction: Heart Disease Present

```
75 sample = pd.DataFrame(  
76     [[63,1,3,145,233,1,0,150,0,2.3,0,0,1]],  
77     columns=X.columns  
78 )  
79 print("The predicted value of the model is: ", model.predict(sample))  
80  
81 prediction = model.predict(sample)  
82 if prediction[0] == 1:  
83     print("Heart Disease Present")  
84 else:  
85     print("No Heart Disease")  
86
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

```
PS E:\Project_Code_Alpha> & "C:\Users\ADSS ABHISHEK\AppData\Local\Programs\Python\Python312\python.exe" e:/Project_Code_Alpha/Disease_prediction.py  
[[73 29]  
[13 90]]  
      precision    recall  f1-score   support  
  
      0         0.85      0.72      0.78         102  
      1         0.76      0.87      0.81         103  
  
 accuracy          0.80  
 macro avg         0.80      0.79      0.79         205  
weighted avg         0.80      0.80      0.79         205  
  
The predicted value of the model is: [1]  
Heart Disease Present  
PS E:\Project Code Alpha>
```

Sample2: Prediction: No Heart Disease

```
75 sample1 = pd.DataFrame(  
76     [[52,1,0,125,212,0,1,168,0,1.0,2,2,3]],  
77     columns=X.columns  
78 )  
79  
80 print("The predicted value of the model is: ", model.predict(sample1))  
81  
82 prediction = model.predict(sample1)  
83 if prediction[0] == 1:  
84     print("Heart Disease Present")  
85 else:  
86     print("No Heart Disease")
```

PROBLEMS OUTPUT DEBUG CONSOLE **TERMINAL** PORTS

PS E:\Project_code_Alpha> & "C:\Users\ADSS ABHISHEK\AppData\Local\Programs\Python\Python312\python.exe" e:/Project_code_Alpha/Disease_prediction.p
[[73 29]
[13 90]]

	precision	recall	f1-score	support
0	0.85	0.72	0.78	102
1	0.76	0.87	0.81	103
accuracy			0.80	205
macro avg	0.80	0.79	0.79	205
weighted avg	0.80	0.80	0.79	205

The predicted value of the model is: [0]
No Heart Disease

Conclusion

This project successfully developed a Heart Disease Prediction System using Logistic Regression. The model was trained on a dataset containing 1025 patient records and achieved an accuracy of approximately 79.5%. The project demonstrates the practical application of machine learning in healthcare and can assist in early disease detection and support clinical decision-making.

References

1. UCI Machine Learning Repository – Heart Disease Dataset
2. Kaggle – Heart Disease Dataset
3. Scikit-Learn Documentation
4. Pandas Documentation
5. Matplotlib Documentation
6. Seaborn Documentation

Using the Random Forest Algorithm: 2nd method Implementation:

Heart Disease Prediction Using Random Forest Algorithm:

Abstract:

Random Forest, a supervised machine learning ensemble classification algorithm, was applied to predict the presence or absence of heart disease based on patient medical data.

Introduction:

Random Forest is an ensemble learning algorithm that combines multiple decision trees and uses majority voting to make predictions. It is widely used in healthcare applications because of its high accuracy and ability to handle complex datasets.

Methodology:

Heart Dataset



Preprocessing



Exploratory Data Analysis



Feature Selection



Train-Test Split



Random Forest Classifier



Prediction



Model Evaluation



Conclusion

Working of Random Forest:

Patient Medical Data



Random Sampling of Data



Build Multiple Decision Trees



Each Tree Predicts:

Disease / No Disease



Majority Voting



Final Prediction

Model Building:

Algorithm Used:

Random Forest Classifier



n_estimators = 100



random_state = 42

Result:

The Random Forest model achieved an accuracy of 98.5%.

The model effectively classified patients with and without heart disease and demonstrated better predictive performance because it combines the decisions of multiple decision trees.





```

Actual values:
[1 1 0 1 0 1 0 0 1 0 1 0 1 1 0 0 0 1 1 0 1 0 0 0 1 1 1 1 0 0]
Predicted values:
[1 1 0 1 0 1 0 0 1 0 1 0 1 1 0 0 0 1 1 0 0 0 0 0 1 1 1 0 0]
Accuracy: 0.9853658536585366
[[102  0]
 [  3 100]]

```

	precision	recall	f1-score	support
0	0.97	1.00	0.99	102
1	1.00	0.97	0.99	103
accuracy			0.99	205
macro avg	0.99	0.99	0.99	205
weighted avg	0.99	0.99	0.99	205

```

Heart Disease Present

```

Conclusion:

This project successfully developed a Heart Disease Prediction System using the Random Forest Classifier. The model was trained on a dataset containing 1025 patient records and achieved an accuracy of 98.5%. The results demonstrate that Random Forest can effectively predict heart disease and can assist healthcare professionals in early disease detection and clinical decision-making.

Algorithm	Accuracy
Logistic Regression	79.5%
Random Forest	98.5%

Random Forest outperformed Logistic Regression because it combines multiple decision trees and can capture complex relationships within medical data more effectively.